

# Hardware Transactional Memory – Why You Should Care

Reversim 2016

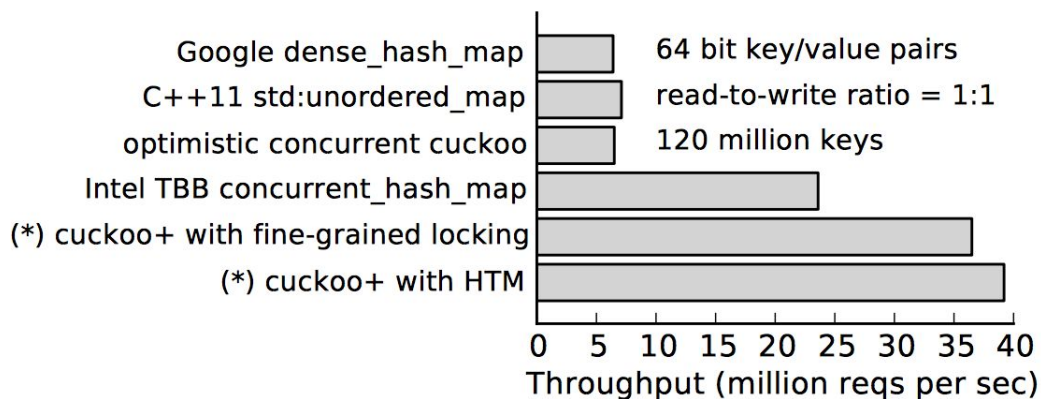
19 Sep 2016

Uri Shamay

Lead Developer, Juno

# Research: Real Cases Improvements

- LevelDB = +25%
- Memcached = x2
- eliminating global interpreter locks in Ruby through HTM = x4.4
- data structures, e.g ::




# Agenda

- concurrency programming evolution
- become familiar with HTM mechanism
- why you should care when you code

```
$>whoami
```

- pretty much... mechanical-sympathy junky
- hard-core infrastructure, scale & high performance solver
- [cmpxchg16.me](http://cmpxchg16.me)
- [@cmpxchg16](https://twitter.com/cmpxchg16)
- [github.com/cmpxchg16](https://github.com/cmpxchg16)

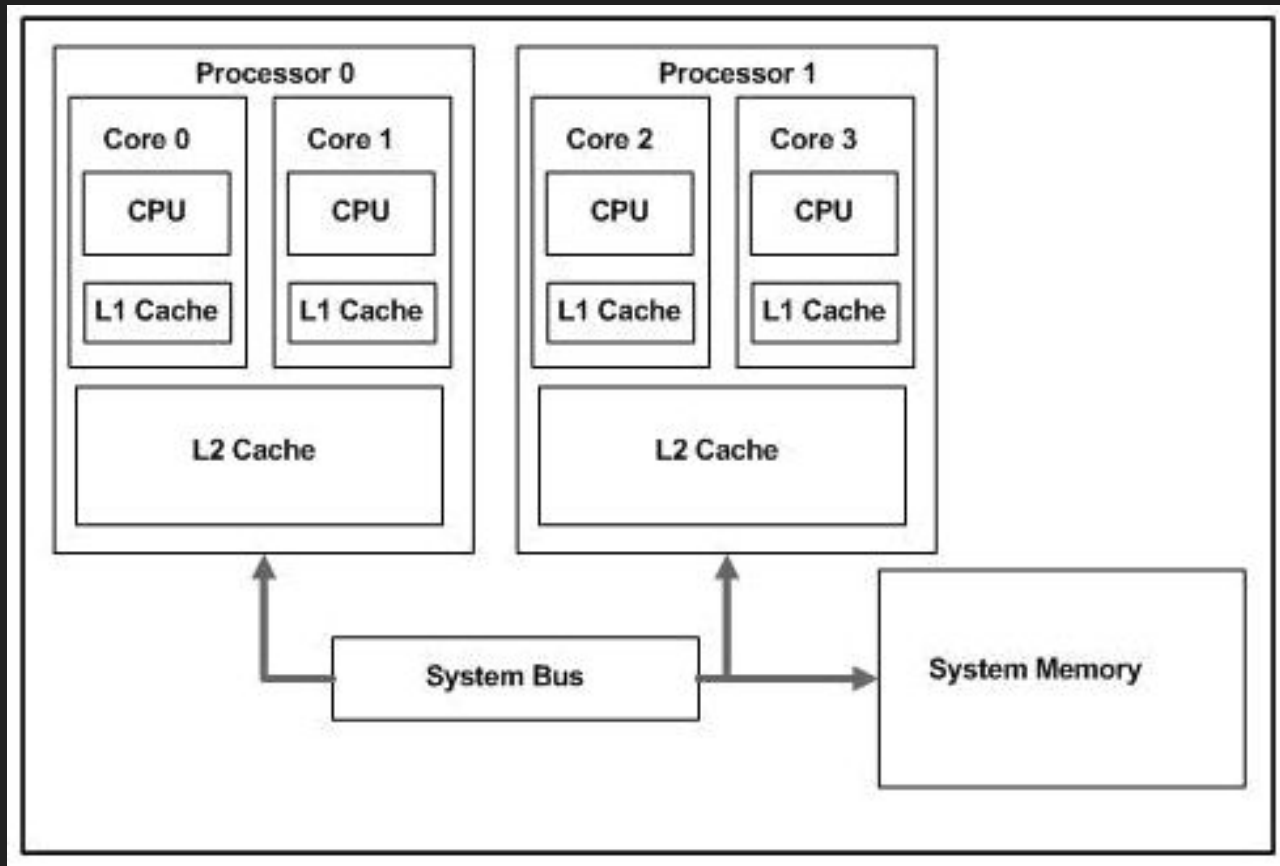
# The First CPU – Intel 4004



[https://en.wikipedia.org/wiki/Intel\\_4004#/media/File:Intel\\_C4004.jpg](https://en.wikipedia.org/wiki/Intel_4004#/media/File:Intel_C4004.jpg)



# Multi-Core CPU



<https://software.intel.com/en-us/articles/software-techniques-for-shared-cache-multi-core-systems>

One CPU Programming – One Thread  
global counter – no inconsistency

```
class Counter {  
    void increment() {  
        this.i++  
    }  
}
```



# Multi-Core CPU Programming

shared counter - using lock

```
class Counter {  
    void synchronized increment() {  
        this.i++  
    }  
}
```

# Multi-Core CPU Programming - hardware-based Optimization

```
class Counter {  
    //shared counter - HW atomic add  
    void increment() {  
        // e.g Java AtomicInteger  
        this.i.incrementAndGet()  
    }  
}
```

# Multi-Core CPU Programming

2 shared counters

```
class Counter {  
    void increment() {  
        // e.g Java AtomicInteger  
        this.counter1.incrementAndGet()  
        this.counter2.incrementAndGet()  
    }  
}
```

but if we want both atomically.. ... ..

# Multi-Core CPU Programming

2 shared counters – what we want (semantically)

```
atomic {  
    this.counter1++  
    this.counter2++  
}
```

# Multi-Core CPU Programming

2 shared counters – how do we get it 0\_\*

```
void synchronized increment() {  
    this.counter1++  
    this.counter2++  
}
```


side note: 2CAS (cmpxchg16b) exists at HW,  
available just in low-level languages

Multi-Core CPU Programming  
and sometimes it becomes more complex...

```
void synchronized mutate(k) {  
    this.hashmap1.remove(k)  
    this.hashmap2.remove(f(k))  
    this.size--  
}
```

# Amdahl's Law

more processors != speedup (mostly..)



[https://en.wikipedia.org/wiki/Amdahl%27s\\_law#/media/File:AmdahlsLaw.svg](https://en.wikipedia.org/wiki/Amdahl%27s_law#/media/File:AmdahlsLaw.svg)

# Hardware Transactional Memory

- a new concurrency primitive to read/write memory atomically in HW (transactional)
- commodity in new Intel-based computers (Haswell/Broadwell microarchitecture)
- Java 8u40 support it `-XX:+UseRTMLocking`



# Hardware Transactional Memory

an optimistic approach to pessimistic locking

```
public class HashMap<K,T> ...  
  
private final array[] = ...  
  
public synchronized T get(K key) {  
    return array[key.hash % ...].get(...);  
}
```

this code **automatically** switches to..

# Hardware Transactional Memory

auto lock elision - using transaction

```
public synchronized T get(K key) {  
    for (int i = 0; i < RETRIES; ++i) {  
        xbegin();  
        T t = array[key.hash % ...].get(...);  
        if (xend() == SUCCEED) {  
            return t  
        }  
    }  
  
    //failed on transaction - fallback to lock  
    lock()  
    ...  
    unlock()  
}
```


# Hardware Transactional Memory

under the hood

- system actions == failed transaction
- no conflict == commit
- conflict == abort transaction & retry
- too many retries == fallback to lock

# Multi-Core CPU – Memory Access

transaction limit = number of lines in L1 cache



# Hardware Transactional Memory

## data contention

```
private int size;
```

```
public synchronized void put(K key, T val) {  
    // different buckets == no data contention  
    array[key.hash % ...].append(val);  
    ++size; // data contention  
}
```

```
public synchronized int size() {  
    return this.size;  
}
```

# Hardware Transactional Memory

reduce data contention

```
private final sizes[] = ...;
```

```
public synchronized void put(K key, T val) {  
    // different buckets == no data contention  
    array[key.hash % ...].append(val);  
    // different buckets == no data contention  
    sizes[key.hash % ...]++;  
}
```

```
public synchronized int size() {  
    for (...) {  
        size += sizes[i];  
    }  
}
```

# Hardware Transactional Memory

## under the hood

- uncontended lock – not interesting
- contended locks with no hotspot gain great improvement
- no conflicts == more parallelization
- data contention should be reduced
- smart implementation (e.g Java), knows to fallback immediate incase of “bad” code

# Hardware Transactional Memory

## wrap-up

- split locks increase parallelization, but with HTM can be more
- HTM != black magic
- profilers provides statistics on HTM, Java:
  - XX:+PrintPreciseRTMLockingStatistics
- HTM usage continues growing



# Hardware Transactional Memory

getting started – check cpu support it

## Option 1:

goto <http://ark.intel.com> && enter the cpu model

look for: **TSX-NI: Yes**

## Option 2:

```
$>lscpu | grep rtm
```

```
Flags:                ... fsgsbase bmi1 hle avx2  
smep bmi2 erms rtm xsaveopt
```

# Thank You!

Uri Shamay

Lead Developer, Juno

**Standing on the shoulders of giants**